

```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

Mounted at /content/drive

+ Code

+ Text

```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

```
data = '/content/drive/MyDrive/Data/weather.csv'
df = pd.read_csv(data)
```

```
df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindSp
0	11/1/2007	Canberra	8.0	24.3	0.0	3.4	6.3	NW	13.000000
1	11/2/2007	Canberra	14.0	26.9	3.6	4.4	9.7	ENE	13.000000
2	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW	13.000000
3	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW	13.000000
4	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE	13.000000

```
df.describe()
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	WindSp
count	366.000000	366.000000	366.000000	366.000000	363.000000	364.000000	359.
mean	7.265574	20.550273	1.428415	4.521858	7.909366	39.840659	9.
std	6.025800	6.690516	4.225800	2.669383	3.481517	13.059807	7.
min	-5.300000	7.600000	0.000000	0.200000	0.000000	13.000000	0.
25%	2.300000	15.025000	0.000000	2.200000	5.950000	31.000000	6.
50%	7.450000	19.650000	0.000000	4.200000	8.600000	39.000000	7.
75%	12.500000	25.500000	0.200000	6.400000	10.500000	46.000000	13.
max	20.900000	35.800000	39.800000	13.800000	13.600000	98.000000	41.

```
df.size
```

```
8784
```

```
df.shape
```

```
(366, 24)
```

```
col_names = df.columns
```

```
col_names
```

```
Index(['Date', 'Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
       'Sunshine', 'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
       'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
       'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',
       'Temp3pm', 'RainToday', 'RISK_MM', 'RainTomorrow'],
      dtype='object')
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 366 entries, 0 to 365
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  366 non-null   object
1   Location              366 non-null   object
2   MinTemp               366 non-null   float64
3   MaxTemp               366 non-null   float64
4   Rainfall              366 non-null   float64
5   Evaporation           366 non-null   float64
6   Sunshine              363 non-null   float64
7   WindGustDir           363 non-null   object
8   WindGustSpeed         364 non-null   float64
9   WindDir9am            335 non-null   object
10  WindDir3pm            365 non-null   object
11  WindSpeed9am          359 non-null   float64
12  WindSpeed3pm          366 non-null   int64
13  Humidity9am           366 non-null   int64
14  Humidity3pm           366 non-null   int64
15  Pressure9am           366 non-null   float64
16  Pressure3pm           366 non-null   float64
17  Cloud9am              366 non-null   int64
18  Cloud3pm              366 non-null   int64
19  Temp9am               366 non-null   float64
20  Temp3pm               366 non-null   float64
21  RainToday             366 non-null   object
22  RISK_MM               366 non-null   float64
23  RainTomorrow          366 non-null   object
dtypes: float64(12), int64(5), object(7)
memory usage: 68.8+ KB
```

```
categorical = [var for var in df.columns if df[var].dtype=='O']


print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :', categorical)
```

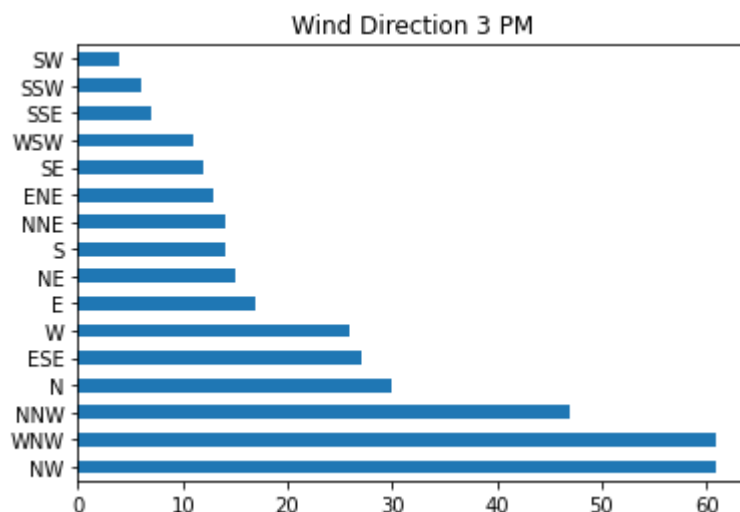
There are 7 categorical variables

The categorical variables are : ['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'WindDir12pm', 'WindDir6pm']

```
df.WindDir9am.value_counts().plot(kind = 'barh')
plt.title("Wind Direction 9 am")
plt.show()
```



```
df.WindDir3pm.value_counts().plot(kind="barh")
plt.title("Wind Direction 3 PM")
plt.show()
```




```
df['RainToday']=df['RainToday'].apply(lambda x:1 if x == "Yes" else 0)
```

```
df['RainTomorrow']=df['RainTomorrow'].apply(lambda x:1 if x == "Yes" else 0)
```

```
df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir
0	11/1/2007	Canberra	8.0	24.3	0.0	3.4	6.3	NW
1	11/2/2007	Canberra	14.0	26.9	3.6	4.4	9.7	ENE
2	11/3/2007	Canberra	13.7	23.4	3.6	5.8	3.3	NW
3	11/4/2007	Canberra	13.3	15.5	39.8	7.2	9.1	NW
4	11/5/2007	Canberra	7.6	16.1	2.8	5.6	10.6	SSE

```
df.WindGustDir.value_counts().plot(kind = "barh",color = 'c')
plt.title("Wind Gust Direction")
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
df=df.dropna()
```

```
df.shape
```

```
(328, 24)
```

```
df.WindGustDir = le.fit_transform(df.WindGustDir)
```

```
df.WindGustDir = le.fit_transform(df.WindGustDir)
```

```
df.WindDir3pm = le.fit_transform(df.WindDir3pm)
```

```
df.WindDir9am = le.fit_transform(df.WindDir9am)
```

```
df.describe()
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGust
count	328.000000	328.000000	328.000000	328.000000	328.000000	328.000000	328.0
mean	7.742988	20.897561	1.440854	4.702439	8.014939	6.192073	40.3
std	5.945199	6.707310	4.289427	2.681183	3.506646	4.337765	13.1
min	-5.300000	7.600000	0.000000	0.200000	0.000000	0.000000	13.0
25%	2.850000	15.500000	0.000000	2.550000	6.000000	2.000000	31.0
50%	7.900000	20.400000	0.000000	4.400000	8.750000	6.500000	39.0
75%	12.800000	25.800000	0.200000	6.600000	10.700000	8.000000	46.0
max	20.900000	35.800000	39.800000	13.800000	13.600000	15.000000	98.0

```
plt.figure(figsize=(15,10))
```

```
plt.subplot(2, 2, 1)
fig = df.boxplot(column='Rainfall')
fig.set_title('')
fig.set_ylabel('Rainfall')
```


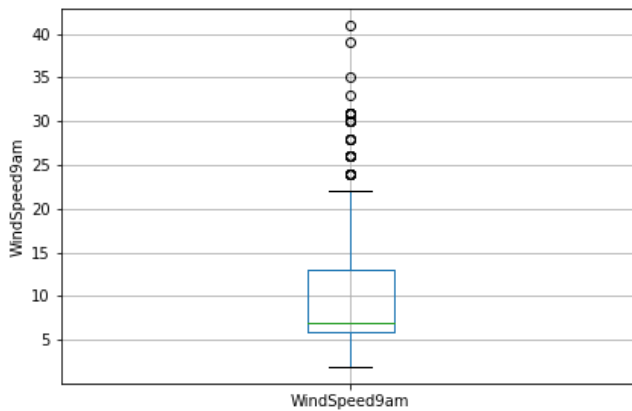
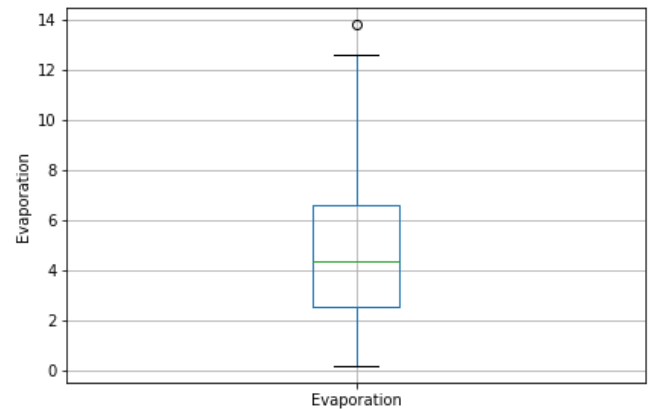
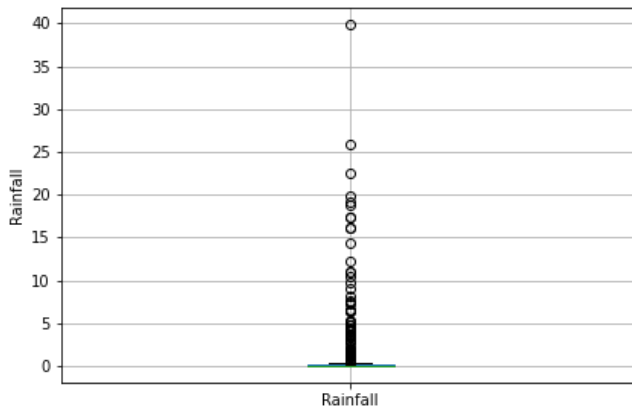
```
plt.subplot(2, 2, 2)
fig = df.boxplot(column='Evaporation')
fig.set_title('')
fig.set_ylabel('Evaporation')
```

```
plt.subplot(2, 2, 3)
fig = df.boxplot(column='WindSpeed9am')
fig.set_title('')
fig.set_ylabel('WindSpeed9am')
```

```
plt.subplot(2, 2, 4)
fig = df.boxplot(column='WindSpeed3pm')
fig.set_title('')
fig.set_ylabel('WindSpeed3pm')
```

```
fig.set_ylabel('WindSpeed3pm')
```

```
Text(0, 0.5, 'WindSpeed3pm')
```



```
plt.figure(figsize=(15,10))
```

```
plt.subplot(2, 2, 1)
fig = df.Rainfall.hist(bins=10)
fig.set_xlabel('Rainfall')
fig.set_ylabel('RainTomorrow')
```





```
plt.subplot(2, 2, 2)
fig = df.Evaporation.hist(bins=10)
fig.set_xlabel('Evaporation')
fig.set_ylabel('RainTomorrow')
```

```
plt.subplot(2, 2, 3)
fig = df.WindSpeed9am.hist(bins=10)
```

```
fig.set_xlabel('WindSpeed9am')
fig.set_ylabel('RainTomorrow')
```

```
plt.subplot(2, 2, 4)
fig = df.WindSpeed3pm.hist(bins=10)
fig.set_xlabel('WindSpeed3pm')
fig.set_ylabel('RainTomorrow')
```

Text(0, 0.5, 'RainTomorrow')



```
IQR = df.Rainfall.quantile(0.75) - df.Rainfall.quantile(0.25)
Lower_fence = df.Rainfall.quantile(0.25) - (IQR * 3)
Upper_fence = df.Rainfall.quantile(0.75) + (IQR * 3)
print('Rainfall outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerbound
```

Rainfall outliers are values < -0.6000000000000001 or > 0.8

```
IQR = df.Evaporation.quantile(0.75) - df.Evaporation.quantile(0.25)
```

```
Lower_fence = df.Evaporation.quantile(0.25) - (IQR * 3)
Upper_fence = df.Evaporation.quantile(0.75) + (IQR * 3)
print('Evaporation outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerbo

Evaporation outliers are values < -9.599999999999996 or > 18.749999999999996
```

```
IQR = df.WindSpeed9am.quantile(0.75) - df.WindSpeed9am.quantile(0.25)
Lower_fence = df.WindSpeed9am.quantile(0.25) - (IQR * 3)
Upper_fence = df.WindSpeed9am.quantile(0.75) + (IQR * 3)
print('WindSpeed9am outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerb

WindSpeed9am outliers are values < -15.0 or > 34.0
```

```
IQR = df.WindSpeed3pm.quantile(0.75) - df.WindSpeed3pm.quantile(0.25)
Lower_fence = df.WindSpeed3pm.quantile(0.25) - (IQR * 3)
Upper_fence = df.WindSpeed3pm.quantile(0.75) + (IQR * 3)
print('WindSpeed3pm outliers are values < {lowerboundary} or > {upperboundary}'.format(lowerb

WindSpeed3pm outliers are values < -28.0 or > 63.0
```

```
from sklearn.model_selection import train_test_split
```

```
x = df.drop(['Date', 'Location', 'RainTomorrow'], axis=1)
```

```
y = df['RainTomorrow']
```

```
train_x , train_y , test_x , test_y = train_test_split(x, y , test_size = 0.2, random_state = 2)
```

```
train_x.shape
```

```
(262, 21)
```

```
train_y.shape
```

```
(66, 21)
```

```
from sklearn.linear_model import LogisticRegression
```

```
model = LogisticRegression()
```

```
model.fit(train_x , test_x)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: Convergence
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```



Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```



```
predict = model.predict(train_y)
```

```
from sklearn.metrics import accuracy_score
```

```
accuracy_score(predict , test_y)
```

```
0.9696969696969697
```

```
from sklearn.svm import SVC
```

```
model1 = SVC()
```

```
model1.fit(train_x , test_x)
```

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

```
predict1 = model1.predict(train_y)
```

```
accuracy_score(predict1 , test_y)
```

```
0.8333333333333334
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
clf = KNeighborsClassifier()
```

```
clf.fit(train_x , test_x)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                     metric_params=None, n_jobs=None, n_neighbors=5, p=2,
                     weights='uniform')
```

```
predict3= clf.predict(train_y)
```

```
accuracy_score(predict3 , test_y)
```

```
0.8939393939393939
```

```
clf1 = KNeighborsClassifier(n_neighbors=4)
```

```
clf1.fit(train_x , test_x)
```

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                     metric_params=None, n_jobs=None, n_neighbors=4, p=2,  
                     weights='uniform')
```

```
predict4 = clf.predict(train_y)
```

```
accuracy_score(predict4 , test_y)
```

```
0.8939393939393939
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf = RandomForestClassifier(max_depth=4)
```

```
rf.fit(train_x , test_x)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,  
                       criterion='gini', max_depth=4, max_features='auto',  
                       max_leaf_nodes=None, max_samples=None,  
                       min_impurity_decrease=0.0, min_impurity_split=None,  
                       min_samples_leaf=1, min_samples_split=2,  
                       min_weight_fraction_leaf=0.0, n_estimators=100,  
                       n_jobs=None, oob_score=False, random_state=None,  
                       verbose=0, warm_start=False)
```

```
predict5=rf.predict(train_y)
```

```
accuracy_score(predict5,test_y)
```

```
1.0
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
model = DecisionTreeClassifier(criterion = "entropy", splitter = "best")
```

```
model.fit(train_x , test_x)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='entropy',
                        max_depth=None, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')
```

```
predict6 = model.predict(train_y)
```

```
accuracy_score(predict6, test_y)
```

```
1.0
```

```
from sklearn.naive_bayes import GaussianNB
```

```
model = GaussianNB()
```

```
model.fit(train_x , test_x)
```

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

```
predict7 = model.predict(train_y)
```

```
accuracy_score(predict7, test_y)
```

```
0.9848484848484849
```

```
from sklearn.neural_network import MLPClassifier
```

```
model = MLPClassifier(hidden_layer_sizes=(10,),
                       max_iter = 5000,
                       activation = 'logistic',
                       solver = 'sgd',
                       learning_rate_init = 0.001
                       )
```

```
model.fit(train_x , test_x)
```

```
MLPClassifier(activation='logistic', alpha=0.0001, batch_size='auto',
               beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
               hidden_layer_sizes=(10,), learning_rate='constant',
```

```
learning_rate_init=0.001, max_fun=15000, max_iter=5000,  
momentum=0.9, n_iter_no_change=10, nesterovs_momentum=True,  
power_t=0.5, random_state=None, shuffle=True, solver='sgd',  
tol=0.0001, validation_fraction=0.1, verbose=False,  
warm_start=False)
```

```
predict8 = model.predict(train_y)
```

```
accuracy_score(predict8,test_y)
```

```
0.8333333333333334
```

```
from sklearn.cluster import KMeans
```

```
model = KMeans(n_clusters = 3, init = "random", algorithm = "full")
```

```
model.fit(train_x , test_x)
```

```
KMeans(algorithm='full', copy_x=True, init='random', max_iter=300, n_clusters=3,  
        n_init=10, n_jobs=None, precompute_distances='auto', random_state=None,  
        tol=0.0001, verbose=0)
```

```
predict9 = model.predict(train_y)
```

```
accuracy_score(predict9,test_y)
```

```
0.15151515151515152
```

```
from sklearn.linear_model import LinearRegression
```

```
model=LinearRegression()
```

```
model.fit(train_x , test_x)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
predict10 = model.predict(train_y)
```

```
from sklearn.metrics import mean_squared_error  
from math import sqrt  
print(sqrt(mean_squared_error(predict10,test_y)))
```

```
0.25684667532443894
```

```

from sklearn.metrics import confusion_matrix

cm =confusion_matrix(predict,test_y)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])

```

Confusion matrix

```
[[55  2]
 [ 0  9]]
```

True Positives(TP) = 55

True Negatives(TN) = 9

False Positives(FP) = 2


False Negatives(FN) = 0

```
import seaborn as sns
```

```
cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                        index=['Predict Positive:1', 'Predict Negative:0'])
```

```
sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f0a69177350>



```
from sklearn.metrics import classification_report
```

```
print(classification_report(predict,test_y))
```

	precision	recall	f1-score	support
0	1.00	0.96	0.98	57
1	0.82	1.00	0.90	9
accuracy			0.97	66
macro avg	0.91	0.98	0.94	66
weighted avg	0.98	0.97	0.97	66

```
TP = cm[0,0]
```

```
TN = cm[1,1]
```

```
FP = cm[0,1]
```

```
FN = cm[1,0]
```

```
classification_accuracy = (TP + TN) / float(TP + TN + FP + FN)
```

```
print('Classification accuracy : {0:0.4f}'.format(classification_accuracy))
```

```
Classification accuracy : 0.9697
```

```
classification_error = (FP + FN) / float(TP + TN + FP + FN)
```

```
print('Classification error : {0:0.4f}'.format(classification_error))
```

```
Classification error : 0.0303
```

```
precision = TP / float(TP + FP)
```

```
print('Precision : {0:0.4f}'.format(precision))
```

```
Precision : 0.9649
```

```
recall = TP / float(TP + FN)
```

```
print('Recall or Sensitivity : {0:0.4f}'.format(recall))
```

```
Recall or Sensitivity : 1.0000
```

```
true_positive_rate = TP / float(TP + FN)
```

```
print('True Positive Rate : {0:0.4f}'.format(true_positive_rate))
```

True Positive Rate : 1.0000

```
false_positive_rate = FP / float(FP + TN)
```

```
print('False Positive Rate : {0:0.4f}'.format(false_positive_rate))
```

False Positive Rate : 0.1818

```
specificity = TN / (TN + FP)
```

```
print('Specificity : {0:0.4f}'.format(specificity))
```

Specificity : 0.8182

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize = (18,7))
```

```
classifiers = ['LogisticRegression', 'SVC', 'KNeighbors', 'RandomForest', 'DecisionTree', 'Ga  
              'MLP', 'KMeans']
```

```
accuracy = [0.96, 0.83, 0.89, 1.0, 1.0, 0.98, 0.83, 0.46]
```

```
c = ["red", "green", "orange", "pink", "cyan", "blue", "orange", "yellow"]
```

```
plt.bar(classifiers, accuracy, width= 0.8, align='center',color= c, edgecolor = 'red')
```

```
i = 0
```

```
j = 2.0
```

```
for i in range(len(classifiers)):
```

```
    plt.annotate(accuracy[i], (-0.1 + i, accuracy[i] + j))
```

```
plt.title("Bar graph representing the Accuracy of the Classifiers")
```

```
plt.xlabel('CLASSIFIERS')
```

```
plt.ylabel('ACCURACY')
```

```
plt.savefig('Accuracy.png')
```

```
plt.show()
```

